


K4 bits-vs-context (Llama-3.1-8B-Instruct, real-text NIAH duel)



Model: meta-llama/Llama-3.1-8B-Instruct. Solid = skeptical-v2 (primary accounting, used-decoder-columns only); faint = skeptical-v1 as-measured (full-C fp16 decoder charge, every parquet before 2026-07-23). Dashed triangle = skeptical-v2-int8, an ACCOUNTING PROJECTION ONLY pending its own quality gate (int8_decoder_roundtrip quality is not yet measured) -- never a measured result. TQ baseline rows carry no spectral pack and are accounting-mode-invariant. k4_b2.5 x tq_b3 crossover -- v1: between 8192 and 16384 (~13,076); v2: between 8192 and 16384 (~10,236) | k4_b2.5 x tq_k3v2 crossover -- v1: between 32768 and 65536 (~61,820); v2: between 32768 and 65536 (~42,522). docs/2026-07-15-k4-duel-results.md.