



D matrix (budget=2.5, gpt2-scale)



Synthesis-order ladder (fit-side-only arms)



gpt2 scale = mechanism verdict only, not a paper-scale quality claim (Llama fit-side replication pre-registered before any paper claim; see corpus_transfer_verdict.json's gpt2_yellow_flag). All values read verbatim from corpus_transfer_verdict.json (per_budget['2.5']); nothing recomputed here. Left: $D = 1 - \text{win}(\text{fit} \neq \text{eval}) / \text{win}(\text{fit} = \text{eval})$, mean over eval caches; matched fit==eval cells are the D=0 reference, not a measured D. Right: shuf/uni/bi are the fit-side-only synthesis-order arms (shufcode/uniwiki+unicode/biwiki+biocode); the 10% line is the pre-registered rule-(a) recipe-confirmed threshold ($D_{\text{uni}} < 0.10$). Verdict rule: $D < 0.10$ insensitive; $D > 0.25$ domain-sensitive; else as-measured