


Per-rank subspace overlap between corpus-fitted bases (gpt2-scale)



gpt2 scale = mechanism verdict only, not a paper-scale quality claim (Llama fit-side replication pre-registered before any paper claim; see corpus_transfer_verdict.json's gpt2_yellow_flag). overlap.parquet, kind=='overlap' & centered==False (uncentered $E[kk^T]$ fit, the pack's own whitener); mean squared principal cosine between $\text{span}(\text{dec}[:, :r])$ of two corpus fits, mean over layers, at each rank cutoff r in overlap_ranks. 1.0 = identical subspaces; wiki-code/wiki-null/code-null are the three corpus pairs the Task-6 diagnostics cover.