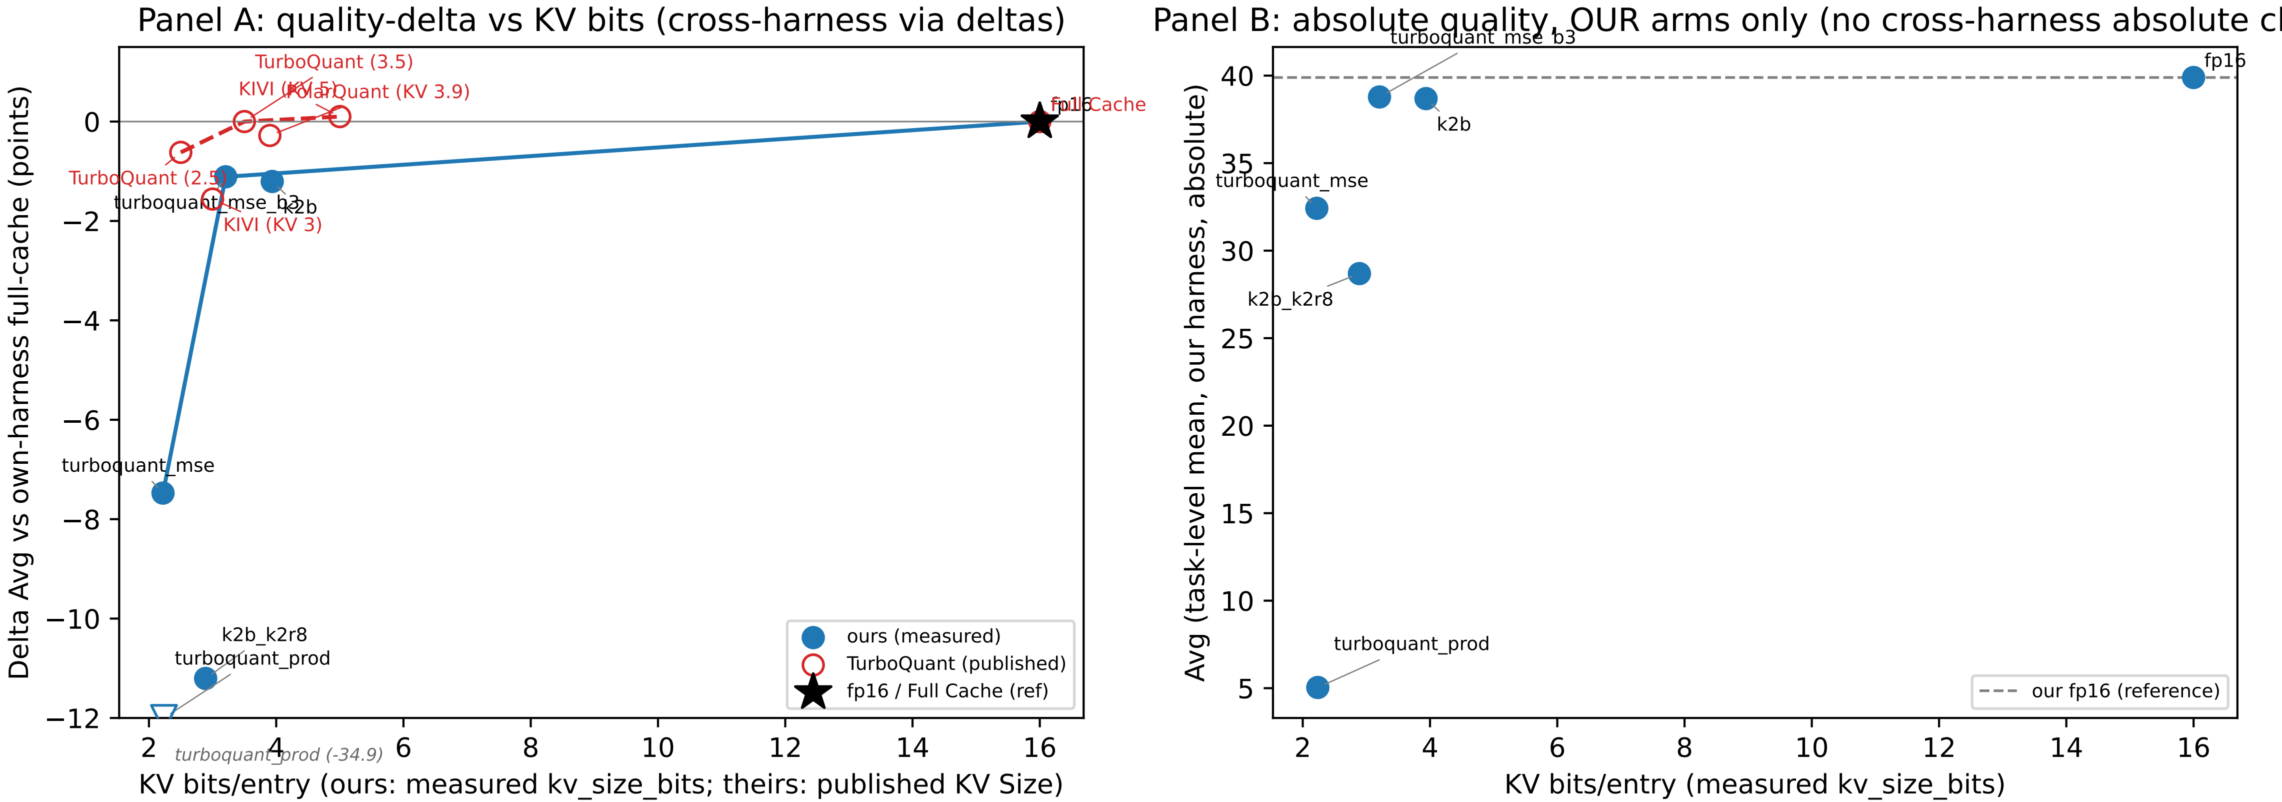

run_id=20260702-164100-46b9579, 20260707-015135-de8b720, 20260707-192442-055dea7. Delta-parity rationale: absolute LongBench Code scores differ ~15pts across harnesses from invisible prompt-policy differences, not quantization (docs/2026-07-06-anchor-forensics-results.md); deltas from each harness's own full-cache Avg cancel that offset and are the only cross-harness-comparable quantity (Panel A). Panel B avoids absolute cross-harness comparison entirely. Bit-width caveat: TurboQuant's 2.5-bit row uses an outlier LongBench split; not an apples-to-apples bit-width match to our measured arms. kivi arm excluded (docs/2026-07-04-kivi-arm-diagnosis.md). Y aggregation is the task-level mean (plot_longbench_table.py's 'Average (all tasks)'), equal weight per task, not the mean of 6 category means. Panel A y-axis clipped to [-12.0, 1.5] for readability; points outside the clip are drawn at the edge with an open arrow marker and their true value is preserved here: turboquant_prod (-34.9).